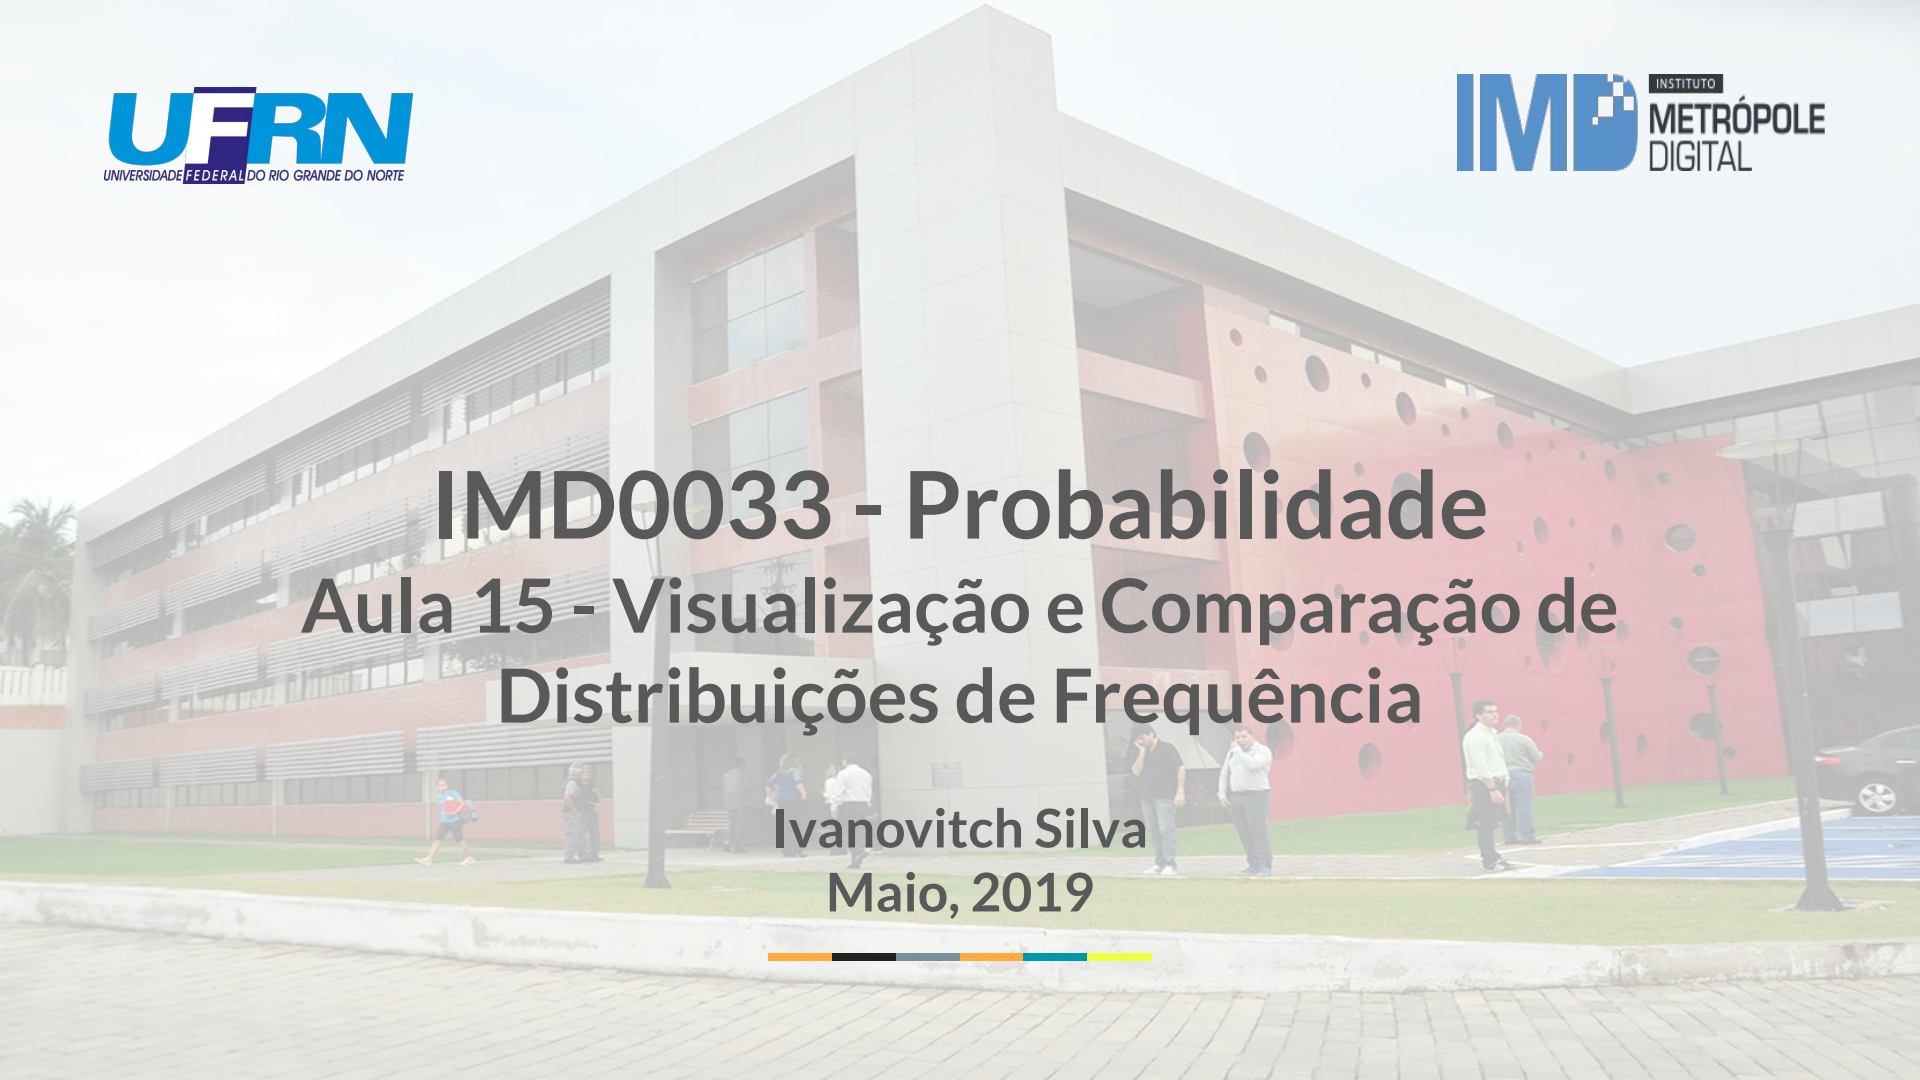

IMD0033 - Probabilidade

Aula 15 - Visualização e Comparação de Distribuições de Frequência

Ivanovitch Silva

Maio, 2019

Agenda (Parte I)

- Visualizando distribuições
- Gráficos de barra, pizza e histogramas
- Assimetria
- Distribuições simétricas

Atualizar o repositório

```
git clone https://github.com/ivanovitchm/imd0033_2019_1.git
```

Ou ....

```
git pull
```

PREVIOUSLY ON...

Data **source**

Getting **good** data

Id	Name	Salary	...	Gender
1	Mary Ann	$35 000	....	Female
2	Marc Downey	$55 000	....	Male
..	....	....	....	....
51	Juliet Ali	$45 000	....	Female
...	...	....	....	....
317	Jane Ace	$95 000	....	Female

Understand how the data is **structured** and **measured**

Organize the data in **comprehensible** forms to find patterns

Gender	Frequency
Male	147
Female	170

Visualize the patterns

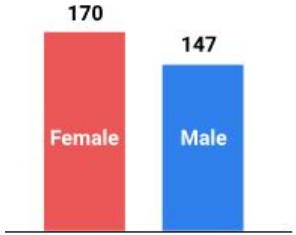

Visualizing Distributions

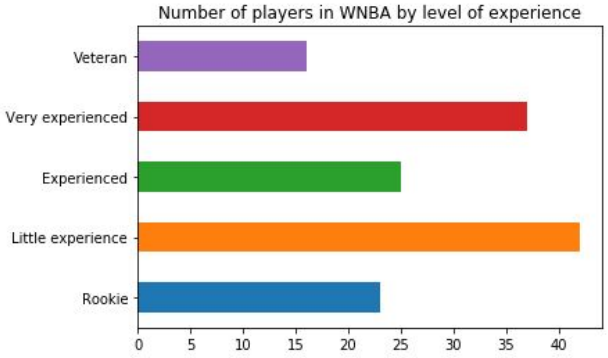

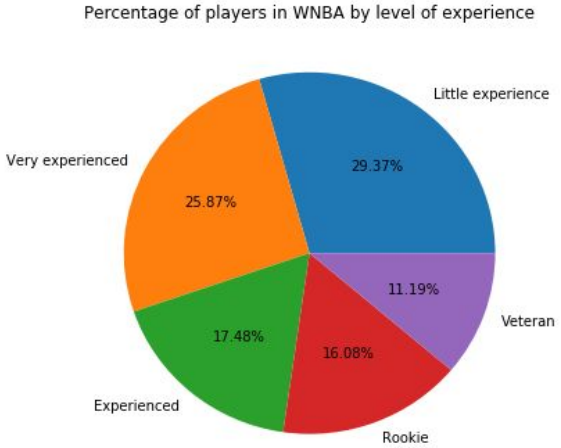

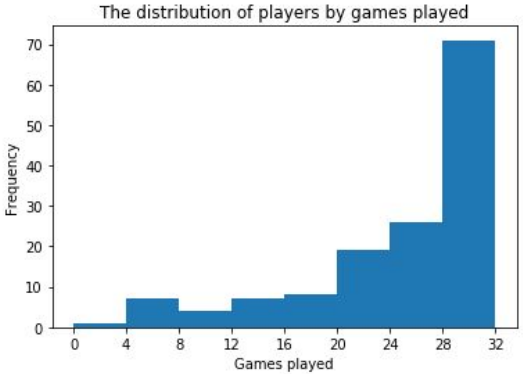

Graphs make easy to scan and compare frequencies, providing us with a single picture of the entire distribution of a variable (**nominal** or **ordinal scale**)

Bar Plots

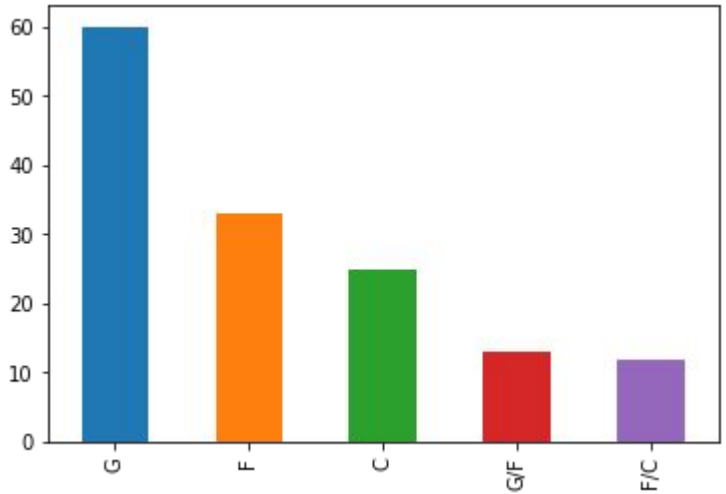


```
wnba['Pos'].value_counts().plot.bar()
```

horizontal bar plots are ideal to use when the labels of the unique values are long

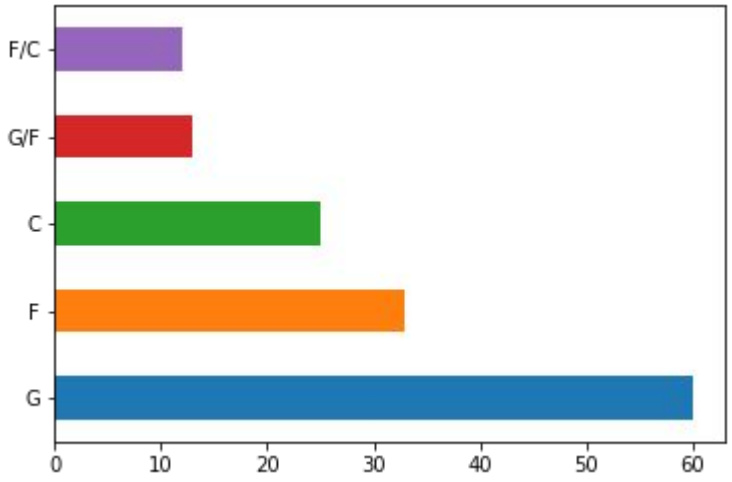


```
wnba['Pos'].value_counts().plot.barh()
```

Pie Charts

Pie Charts

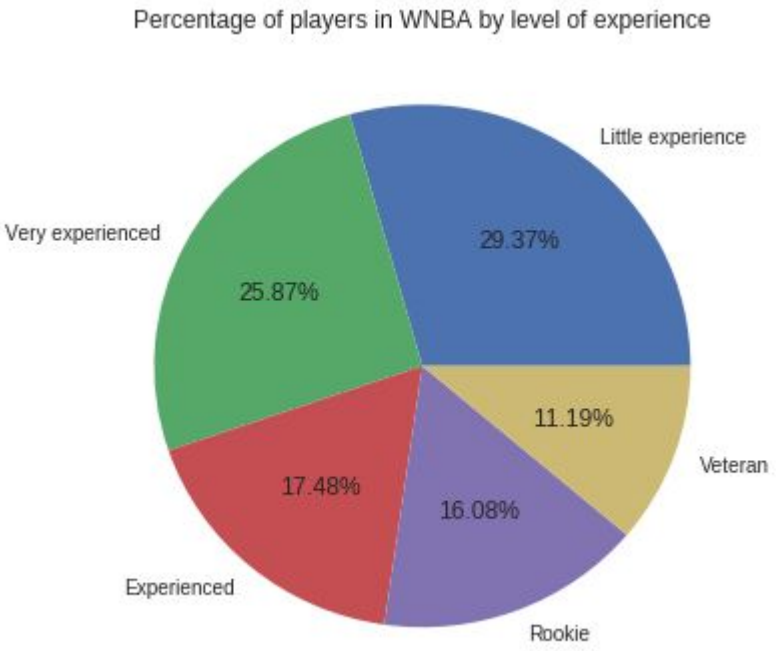


```
wnba['Exp_ordinal'].value_counts().\
plot.pie(figsize = (6,6),
         autopct = '%.2f%%',
         title = 'Percentage of players in \
         WNBA by level of experience')
plt.ylabel('')
```

Histograms

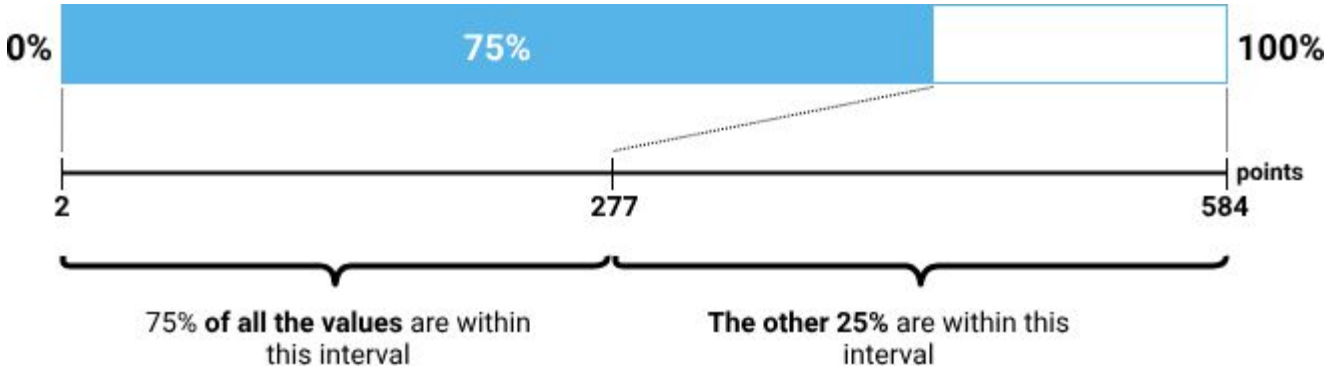

We can see that 75% of the values are distributed within a relatively narrow interval (between 2 and 277), while the remaining 25% are distributed in an interval that's slightly larger.

```
>> wnba['PTS'].describe()
count    143.000000
mean     201.790210
std      153.381548
min        2.000000
25%       75.000000
50%      177.000000
75%      277.500000
max      584.000000
```

The Statistics Behind Histograms

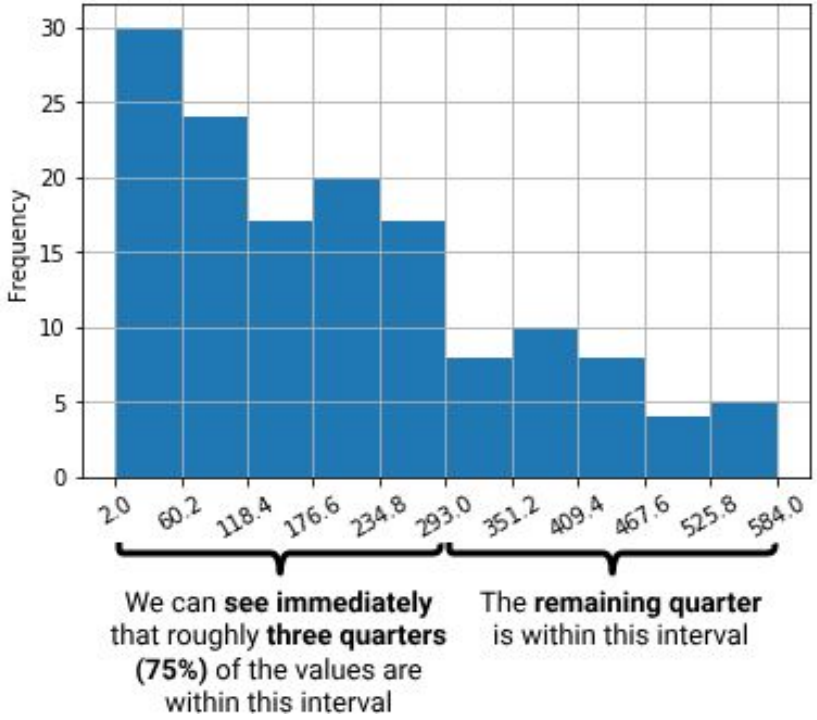


```
>> wnba['PTS'].describe()
count    143.000000
mean     201.790210
std      153.381548
min        2.000000
25%       75.000000
50%      177.000000
75%      277.500000
max      584.000000
Name: PTS, dtype: float64
```

```
>> wnba['PTS'].plot.hist()
```

Binning for Histograms

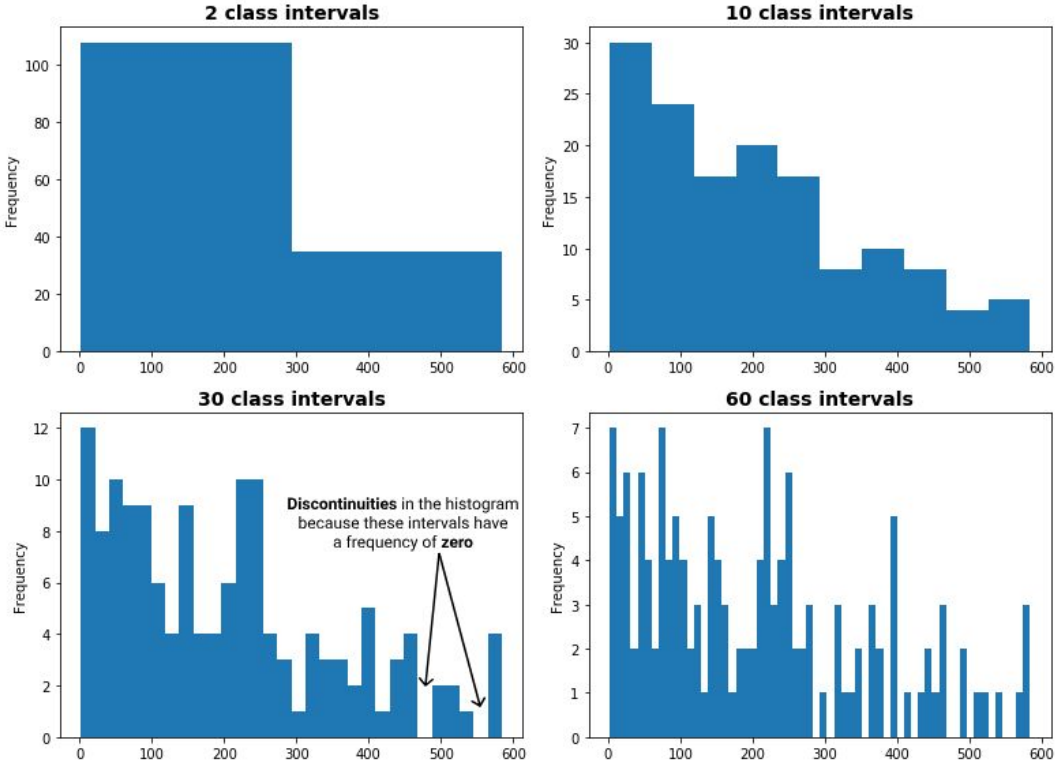

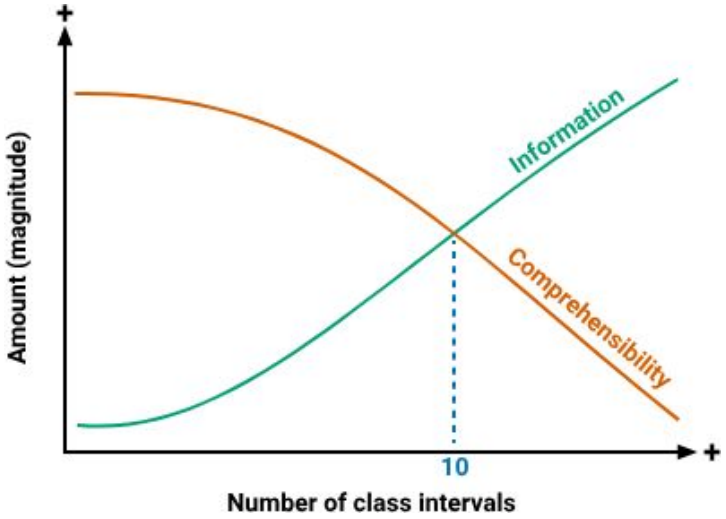

Skewed Distributions

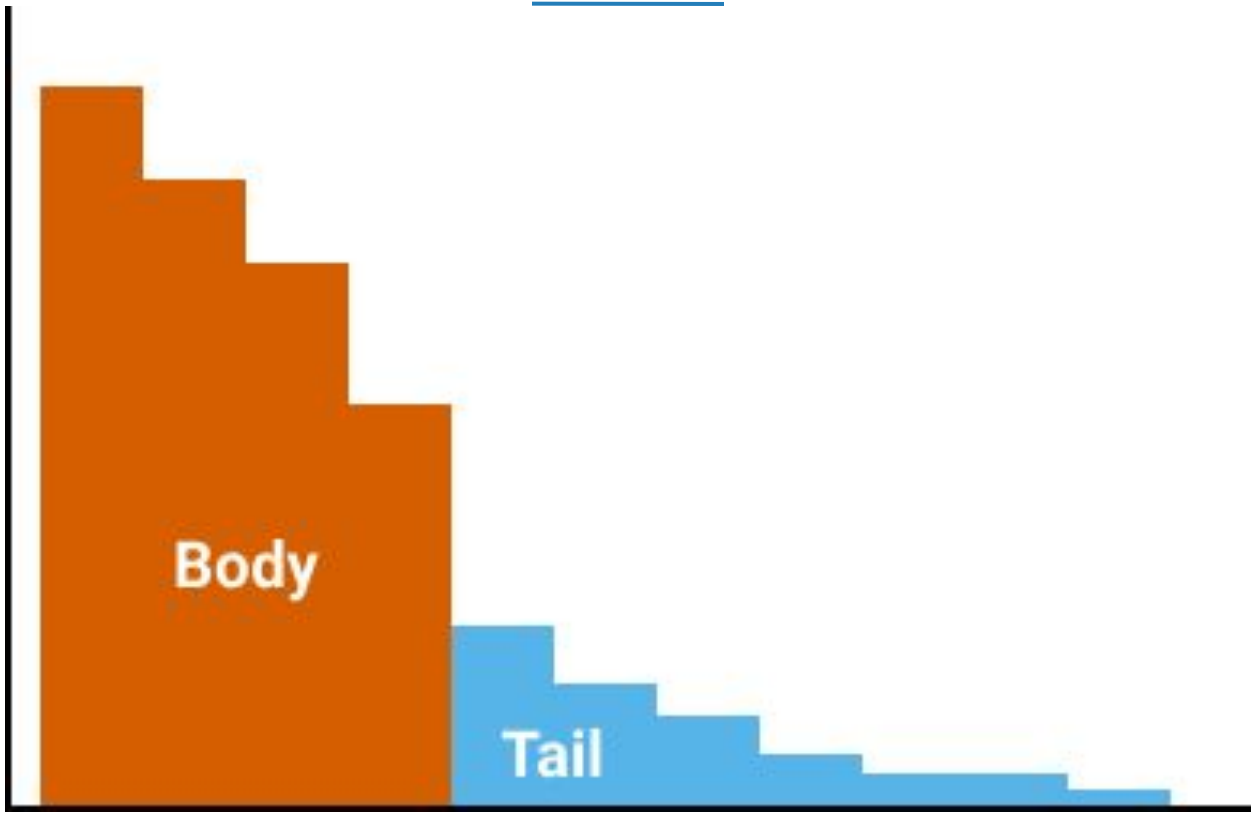

Skewed Distributions

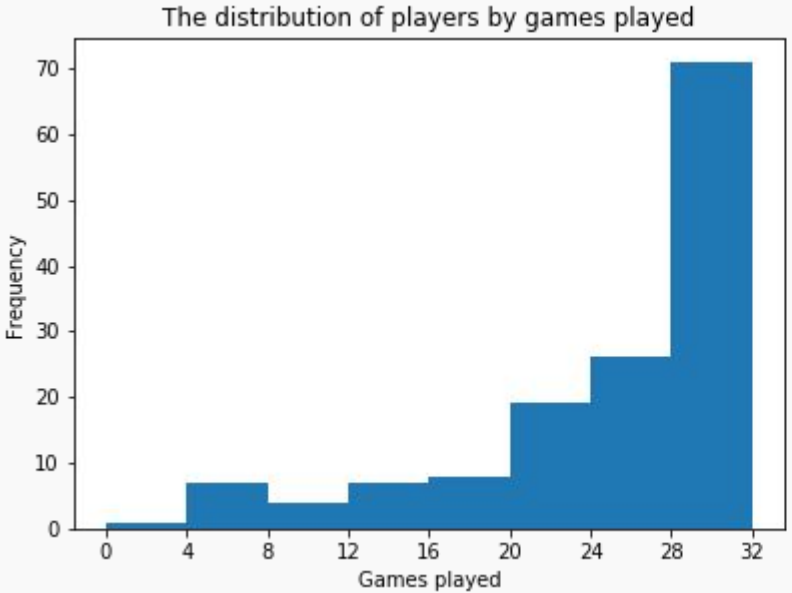

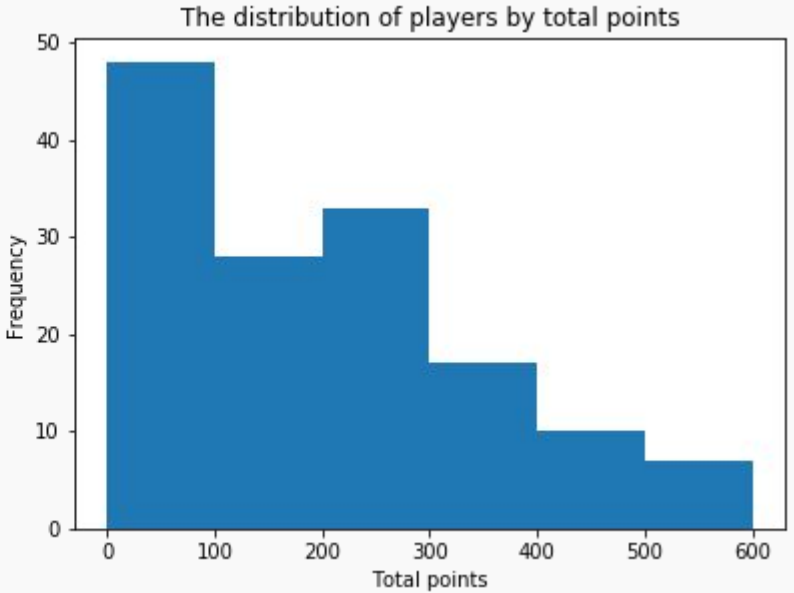

Skewed Distributions

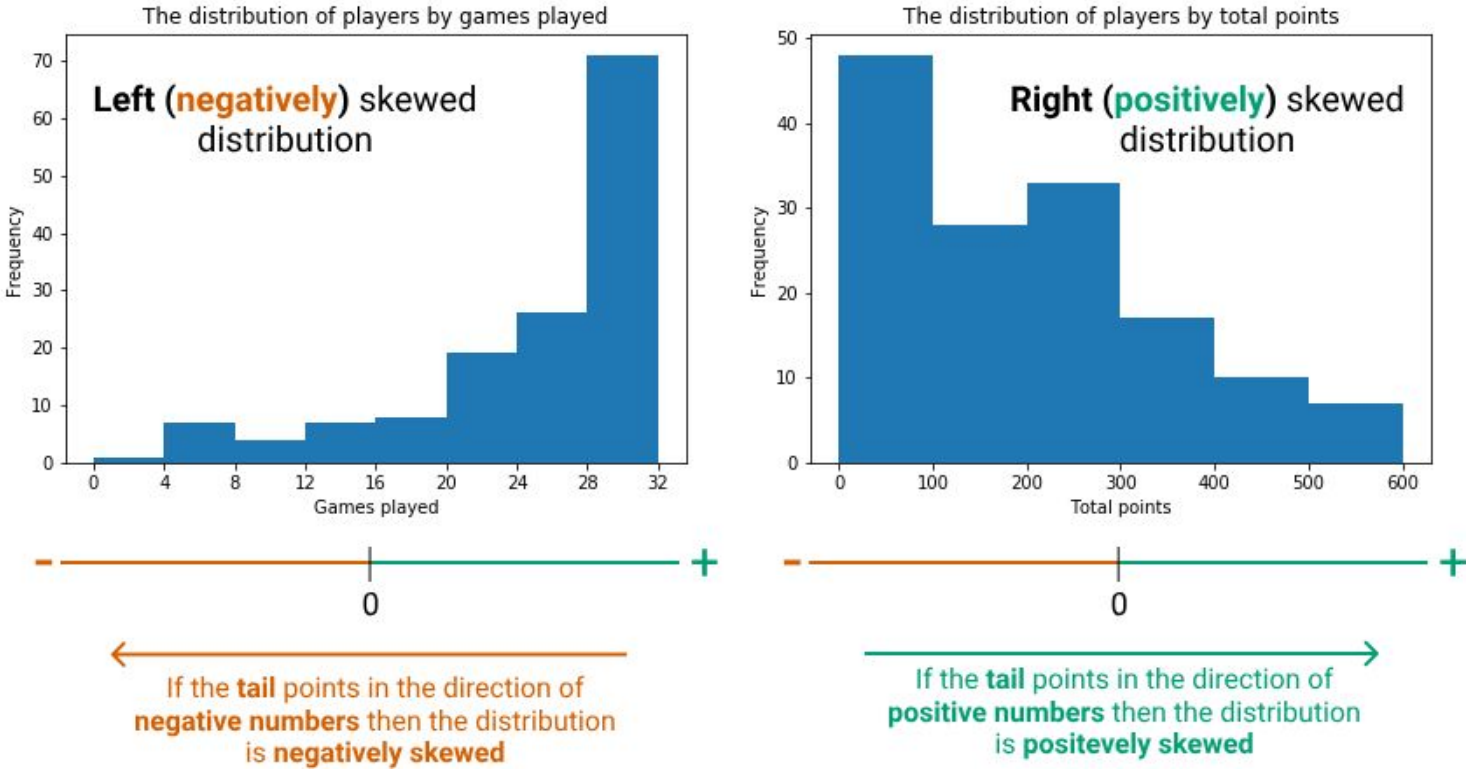

Symmetrical Distributions

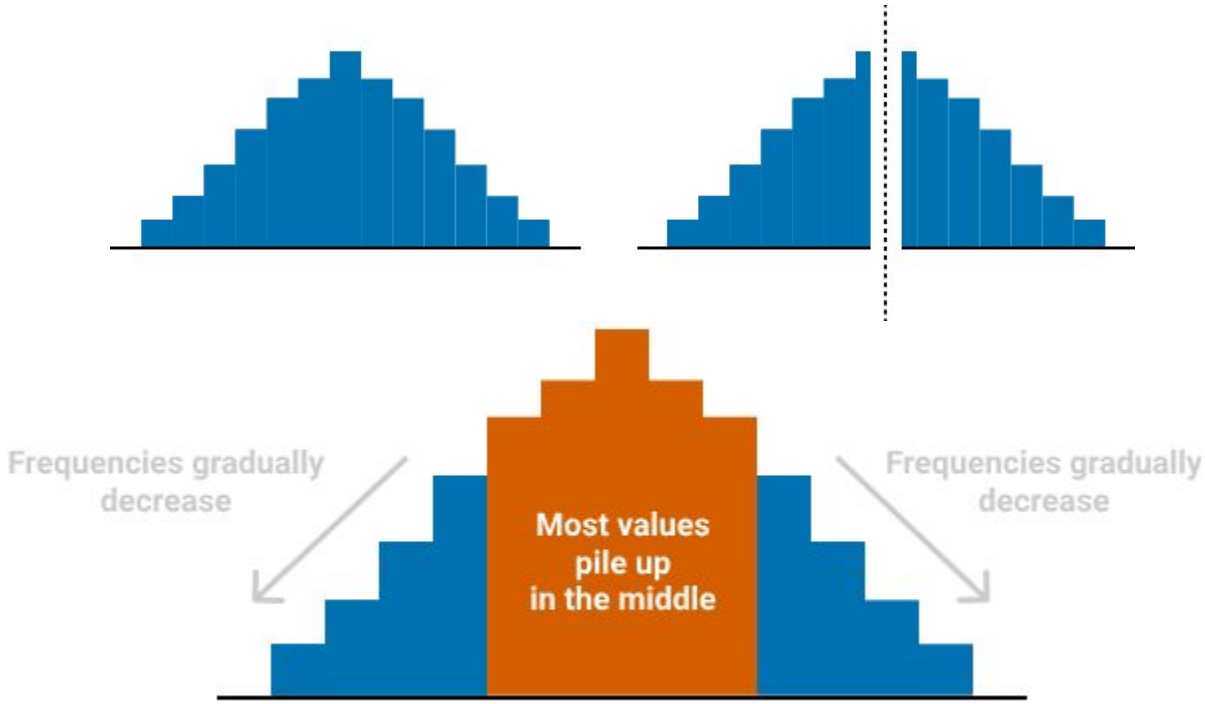

Symmetrical Distribution (uniform)

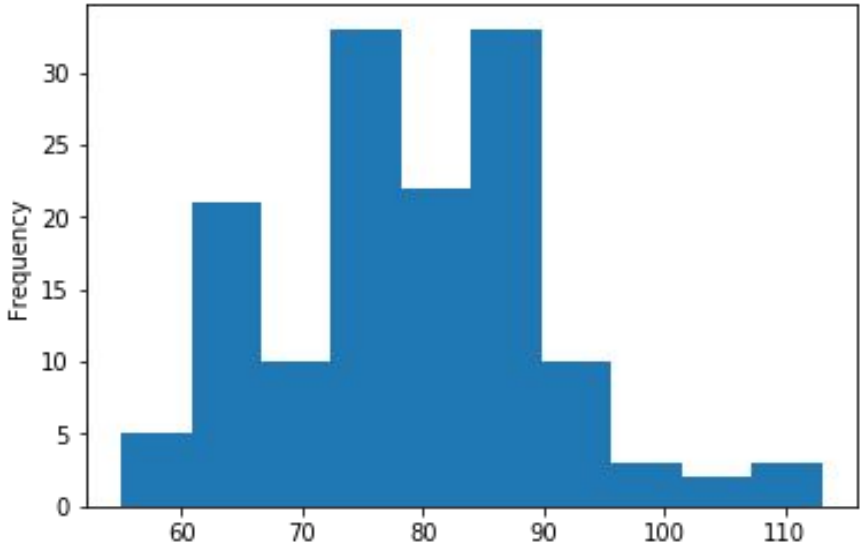

Scale of measurement	Graphs we can use to show the distribution
Nominal	
Ordinal	
Interval	
Ratio	

Data **source**

Getting **good** data

Id	Name	Salary	...	Gender
1	Mary Ann	$35 000	....	Female
2	Marc Downey	$55 000	....	Male
..	....	....	....	....
51	Juliet Ali	$45 000	....	Female
...	...	....	....	....
317	Jane Ace	$95 000	....	Female

Understand how the data is **structured** and **measured**

Organize the data in **comprehensible** forms to find patterns

Gender	Frequency
Male	147
Female	170

Visualize the patterns

Agenda (Parte II)

- Agrupamentos de gráficos de barras
- Comparando histogramas
- Estimativa de densidade kernel
- Gráficos de faixa e caixa
- Pontos fora da curva

Seaborn

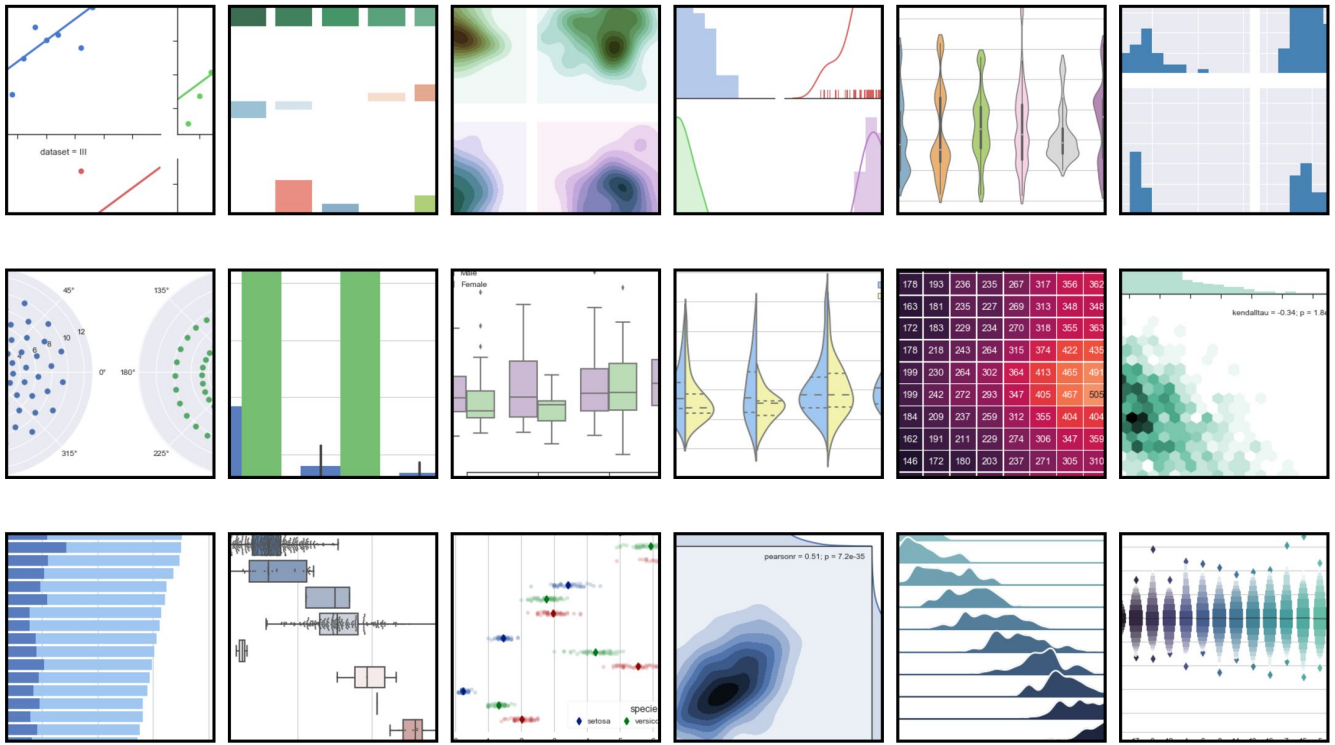
dataset = III
45°
135°
0°
180°
315°
225°
Male
Female
178 193 236 235 267 317 356 362
163 181 235 227 269 313 348 345
172 183 229 234 270 318 355 363
178 218 243 264 315 374 422 435
199 230 264 302 364 413 465 491
199 242 272 293 347 405 467 506
184 209 237 259 312 355 404 404
162 191 211 229 274 306 347 359
146 172 180 203 237 271 305 310
kendalltau = -0.34; p = 1.8
pearsonr = 0.51; p = 7.2e-35
setosa

Comparing Frequency Distribution

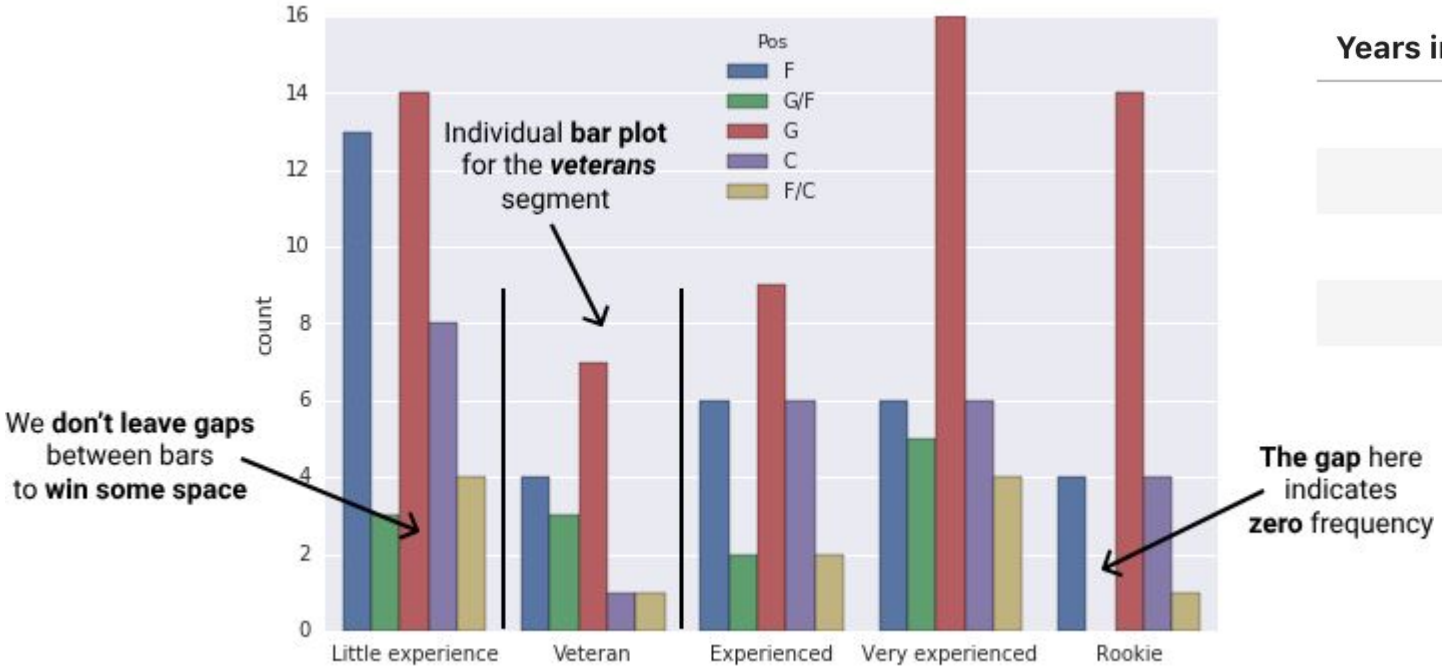

Years in WNBA	Label
0	Rookie
1-3	Little experience
4-5	Experienced
5-10	Very experienced
>10	Veteran

```
sns.countplot(x = 'Exp_ordinal', hue = 'Pos', data = wnba)
```

Challenge: Do older players play less?

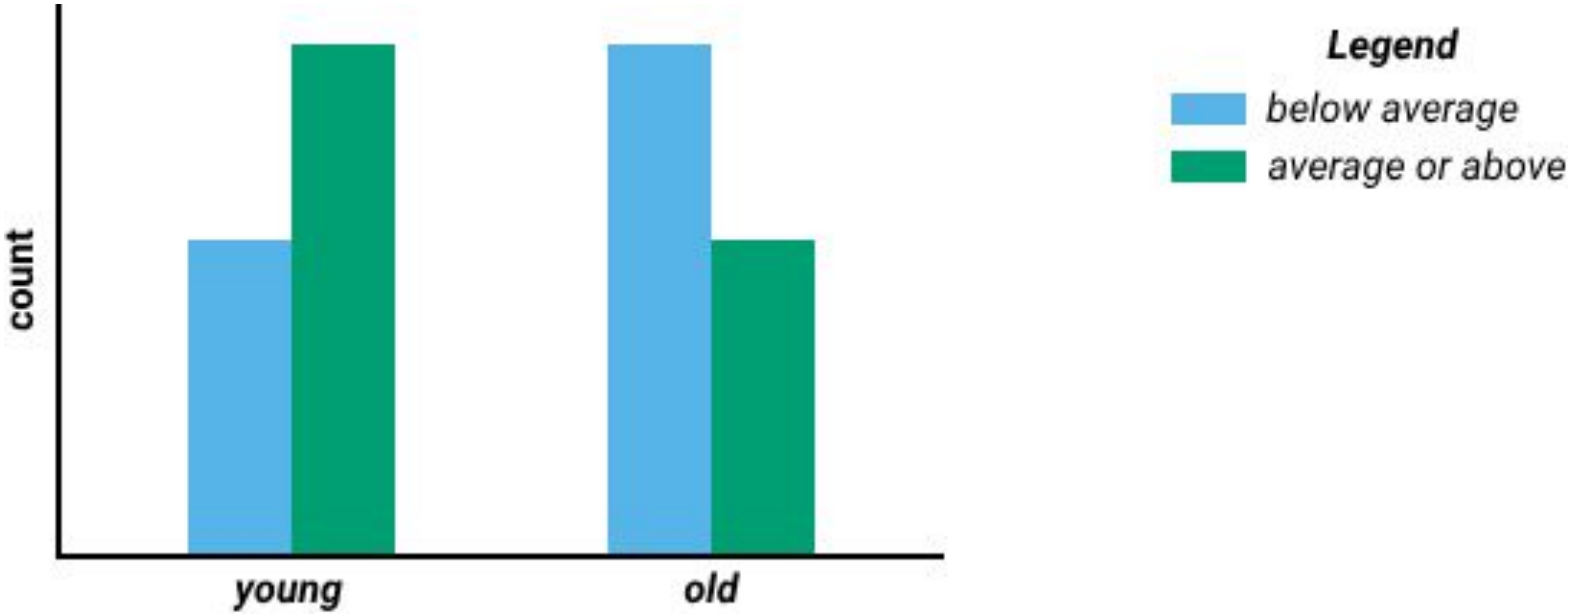

```
wnba['age_mean_relative'] = wnba['Age'].apply(lambda x: 'old' if x >= 27 else 'young')
wnba['min_mean_relative'] = wnba['MIN'].apply(lambda x: 'average or above' if x >= 497 else
                                                'below average')
cols = ["Name","Age","age_mean_relative","MIN","min_mean_relative"]
```

	Name	Age	age_mean_relative	MIN	min_mean_relative
0	Aerial Powers	23	young	173	below average
1	Alana Beard	35	old	947	average or above
2	Alex Bentley	26	young	617	average or above
3	Alex Montgomery	28	old	721	average or above
4	Alexis Jones	23	young	137	below average

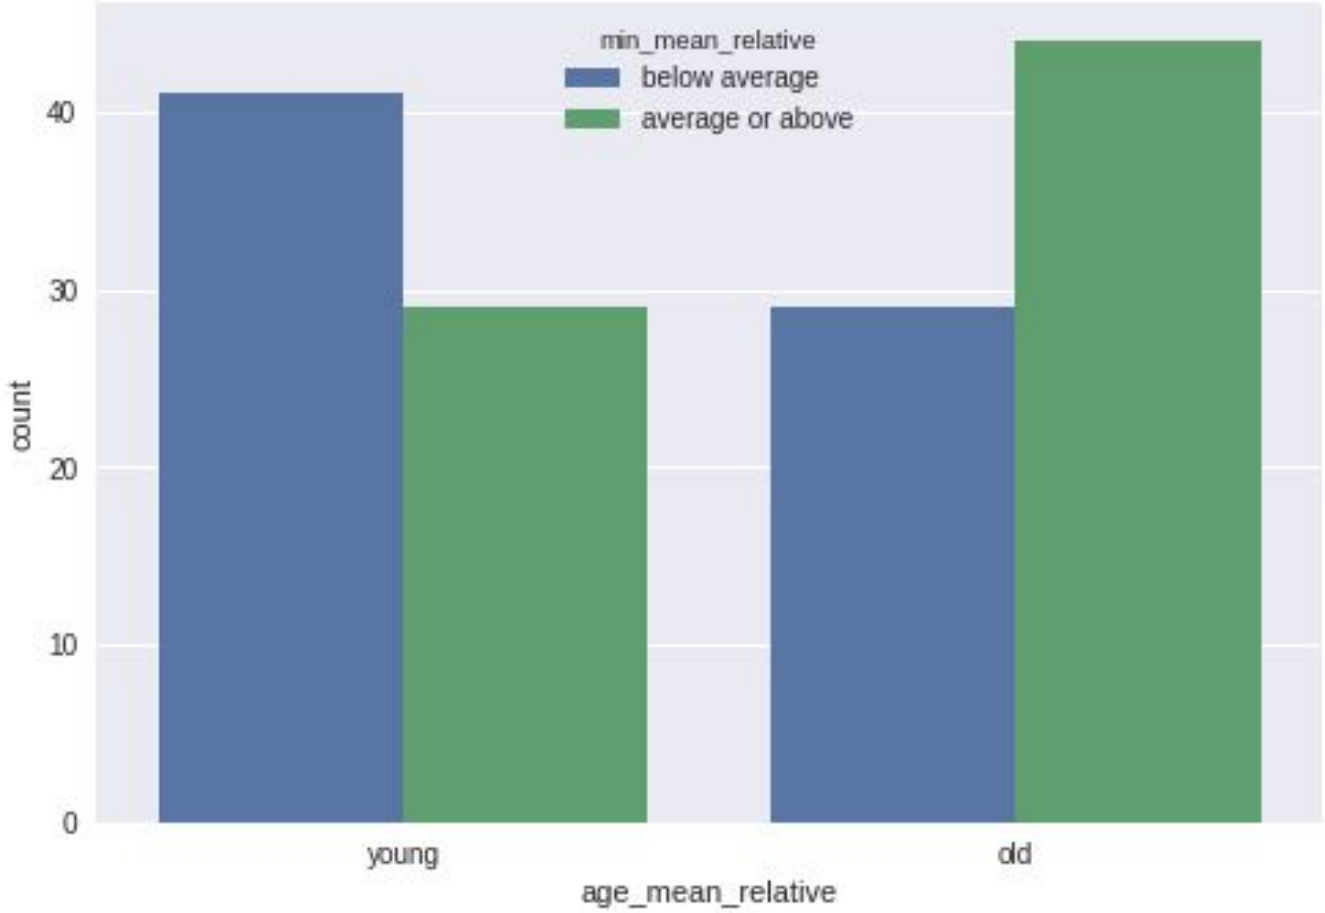


```
sns.countplot(x = 'age_mean_relative', hue = 'min_mean_relative', data = wnba)
```

Comparing Histograms

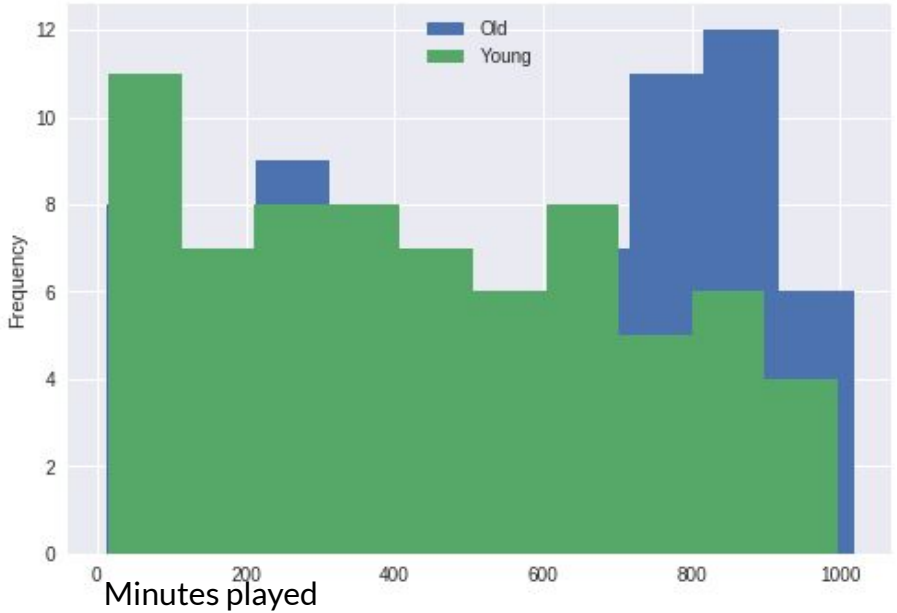


```
wnba[wnba.Age >= 27]['MIN'].plot.hist(label = 'Old', legend = True)
wnba[wnba.Age < 27]['MIN'].plot.hist(label = 'Young', legend = True)
```

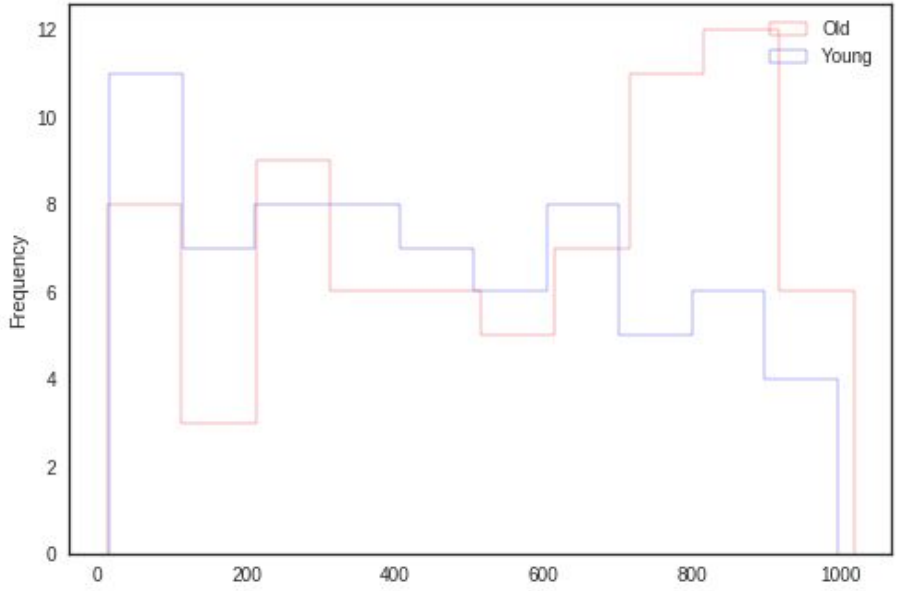


```
sns.set_style("white")
wnba[wnba.Age >= 27]['MIN'].plot.hist(histtype = 'step',
                                      label = 'Old',
                                      legend = True,color="red")
```

Comparing Histograms

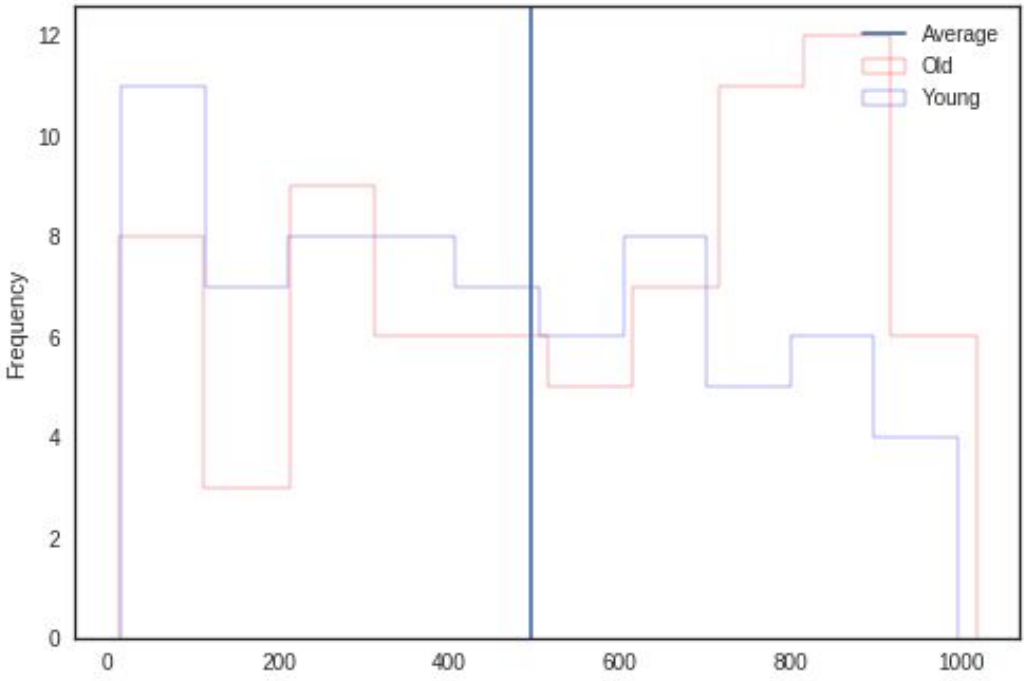


```
plt.axvline(497, label = 'Average')
```

Kernel Density Estimate (KDE) Plots

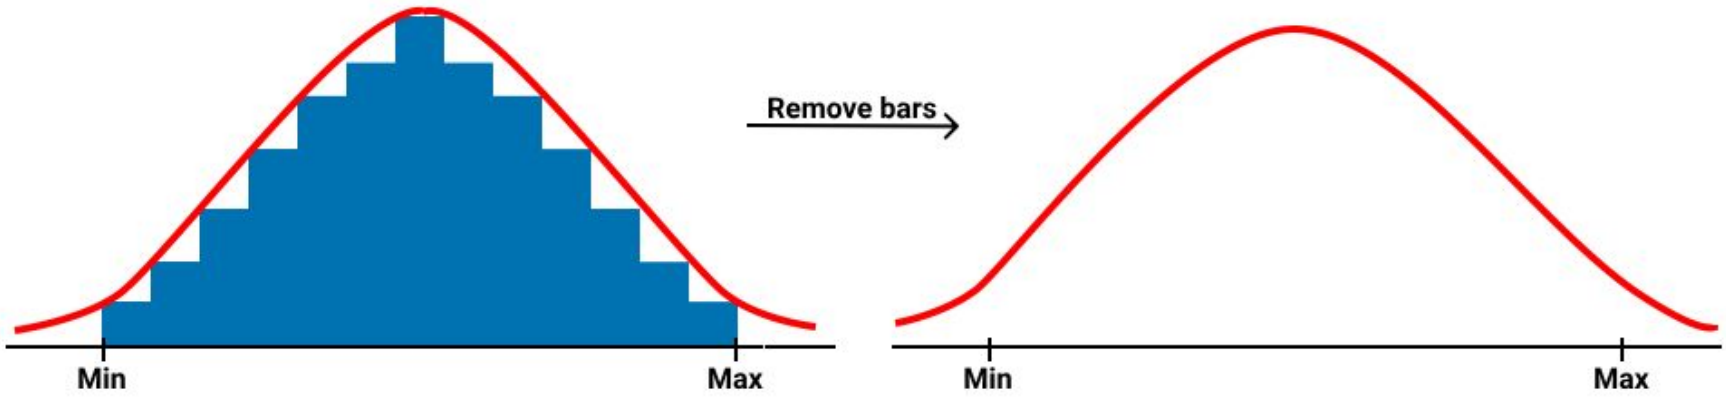

Kernel Density Estimate Plots

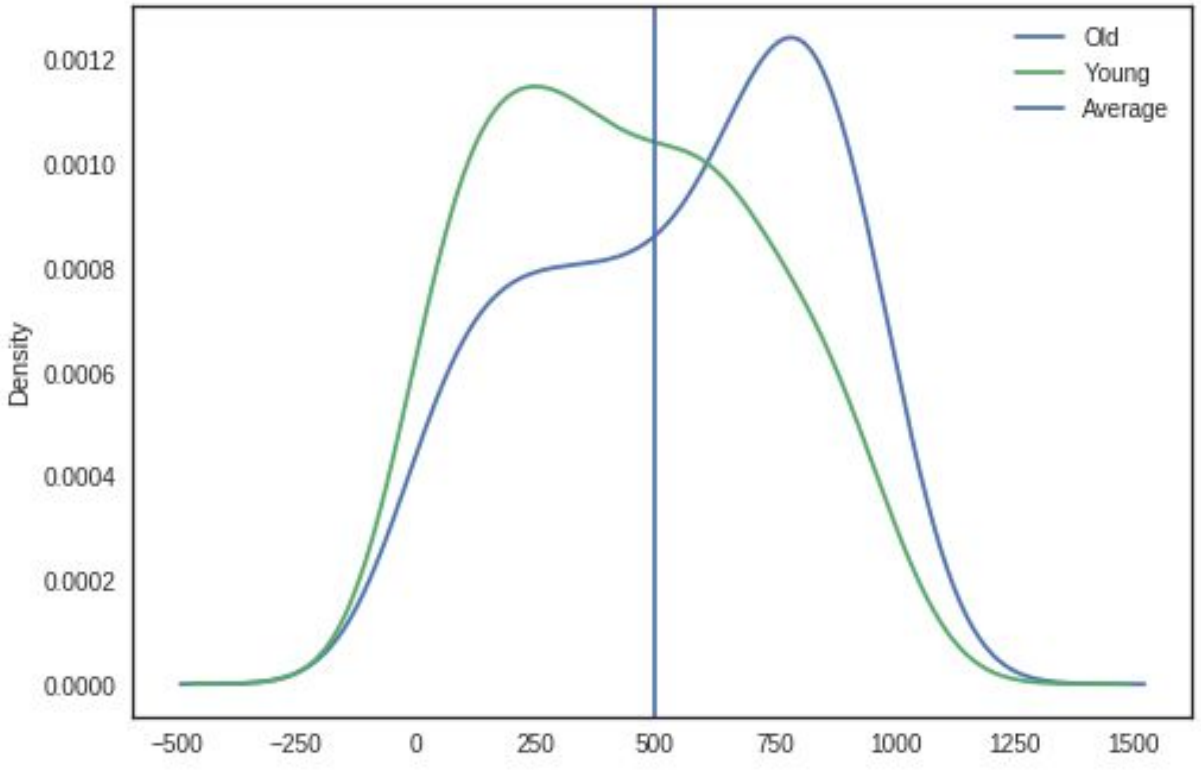

Drawbacks of Kernel Density Plots

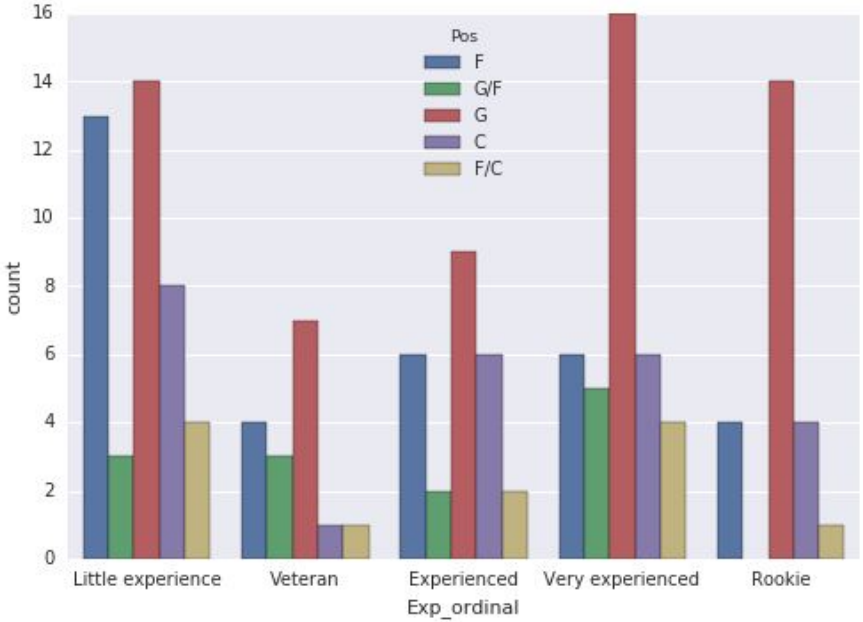

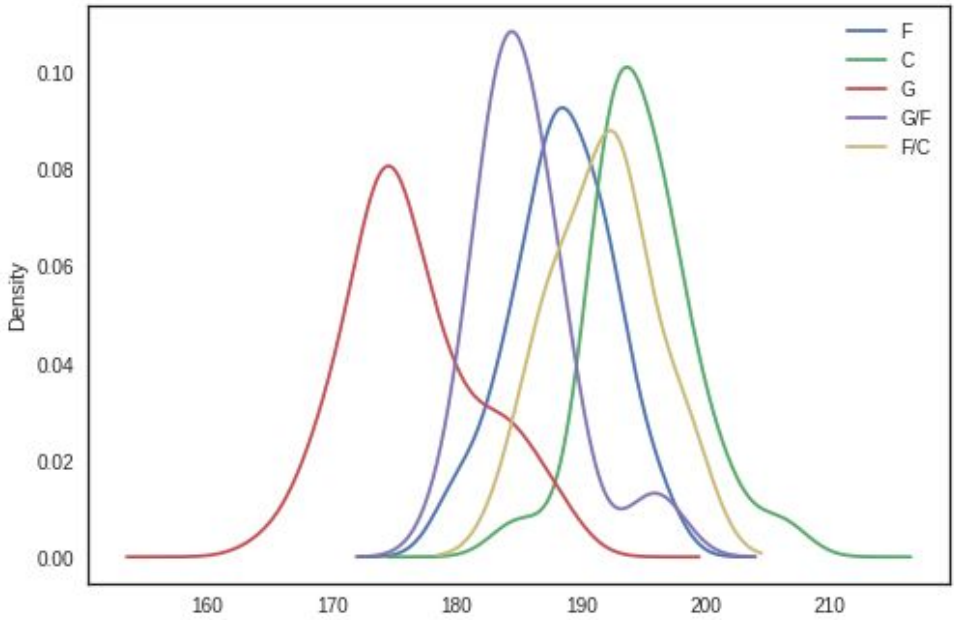

Strip Plots

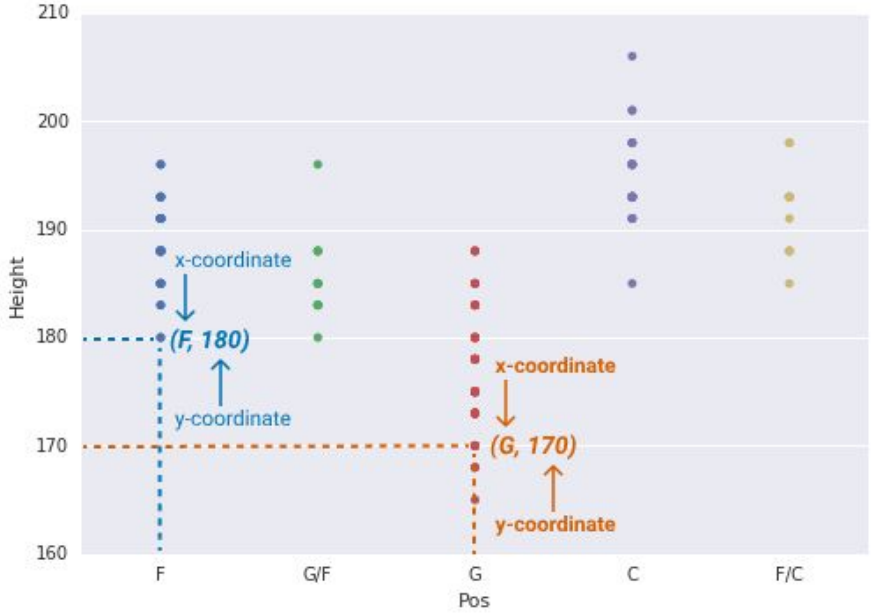
210
200
190
180
170
160
Height
x-coordinate
(F, 180)
y-coordinate
x-coordinate
(G, 170)
y-coordinate
F
G/F
G
C
F/C
Pos

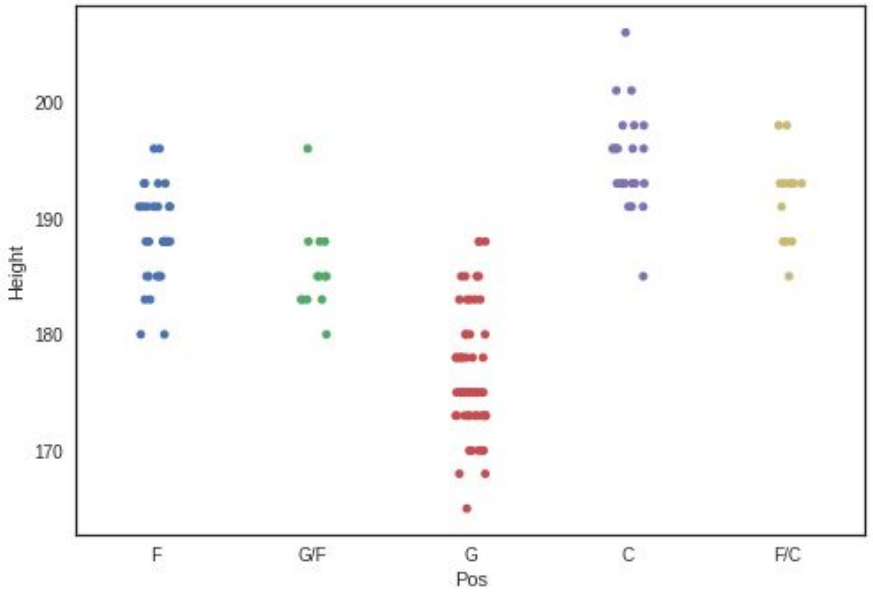
200
190
180
170
Height
F
G/F
G
C
F/C
Pos

Box Plots

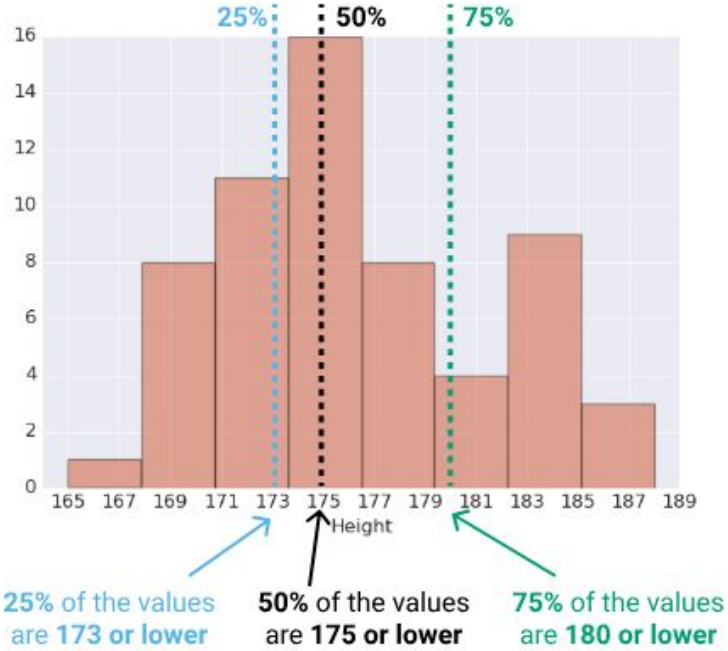
25%
50%
75%
16
14
12
10
8
6
4
2
0
165 167 169 171 173 175 177 179 181 183 185 187 189
Height
25% of the values are 173 or lower
50% of the values are 175 or lower
75% of the values are 180 or lower

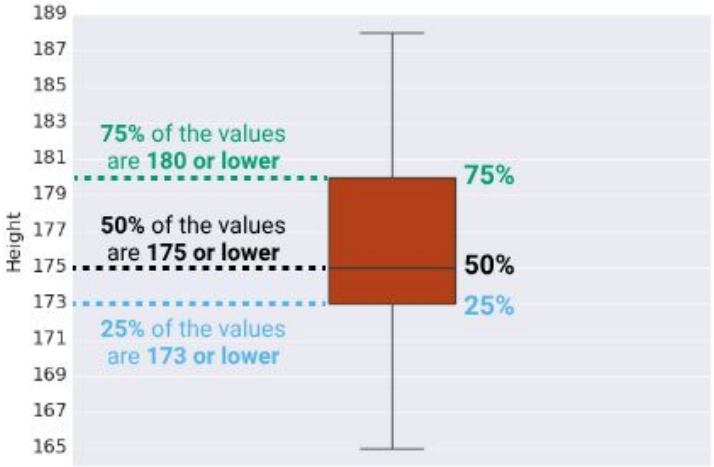
189
187
185
183
181
179
177
175
173
171
169
167
165
Height
75% of the values are 180 or lower
50% of the values are 175 or lower
25% of the values are 173 or lower
75%
50%
25%

Box Plots

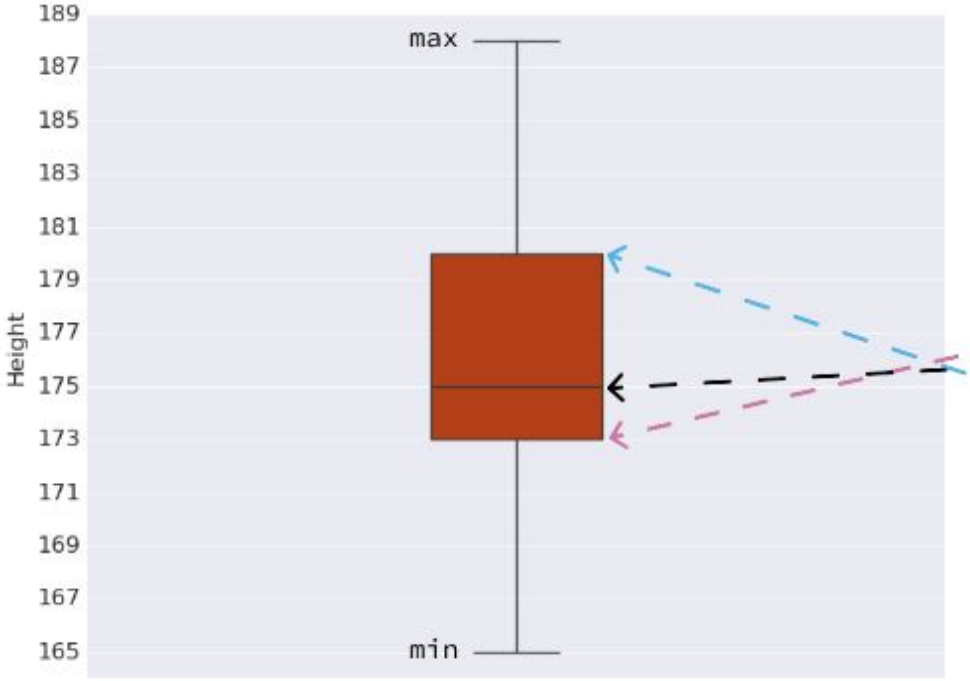


```
>> wnba[wnba['Pos'] == 'G']['Height'].describe()
count     60.000000
mean     176.450000
std        5.312043
min      165.000000
25%      173.000000
50%      175.000000
75%      180.000000
max      188.000000
Name: Height, dtype: float64
```

Outliers

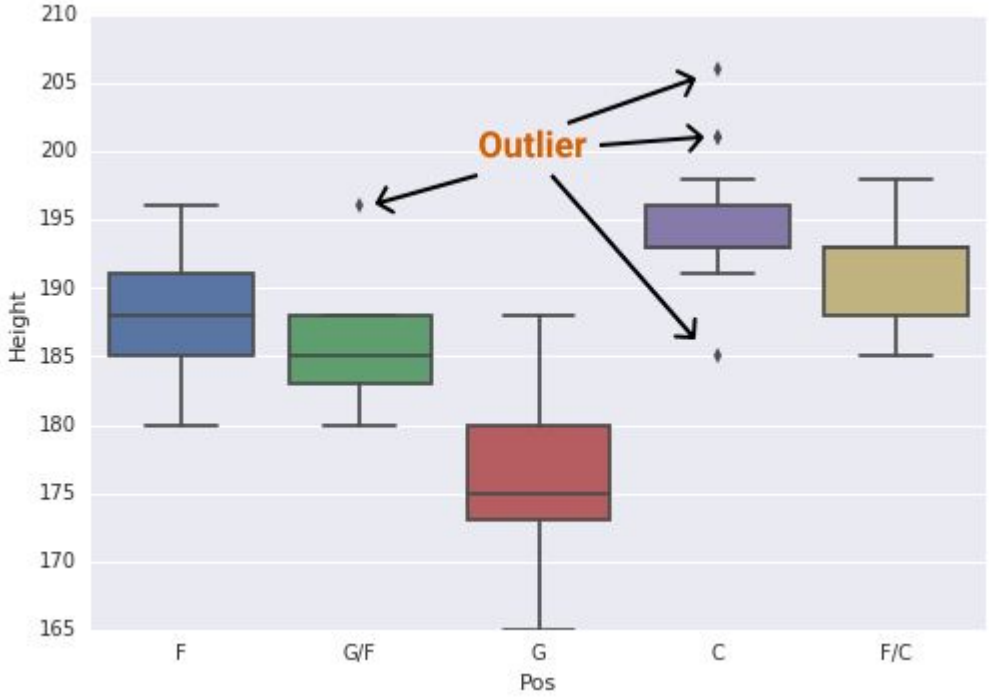
Outlier
Height
210
205
200
195
190
185
180
175
170
165
F
G/F
G
C
F/C
Pos

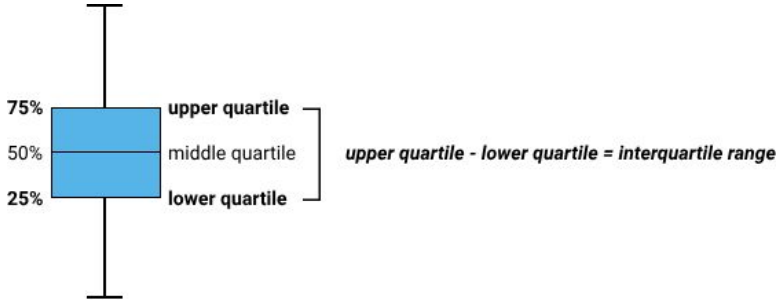
75%
upper quartile
50%
middle quartile
25%
lower quartile
upper quartile - lower quartile = interquartile range

Outliers

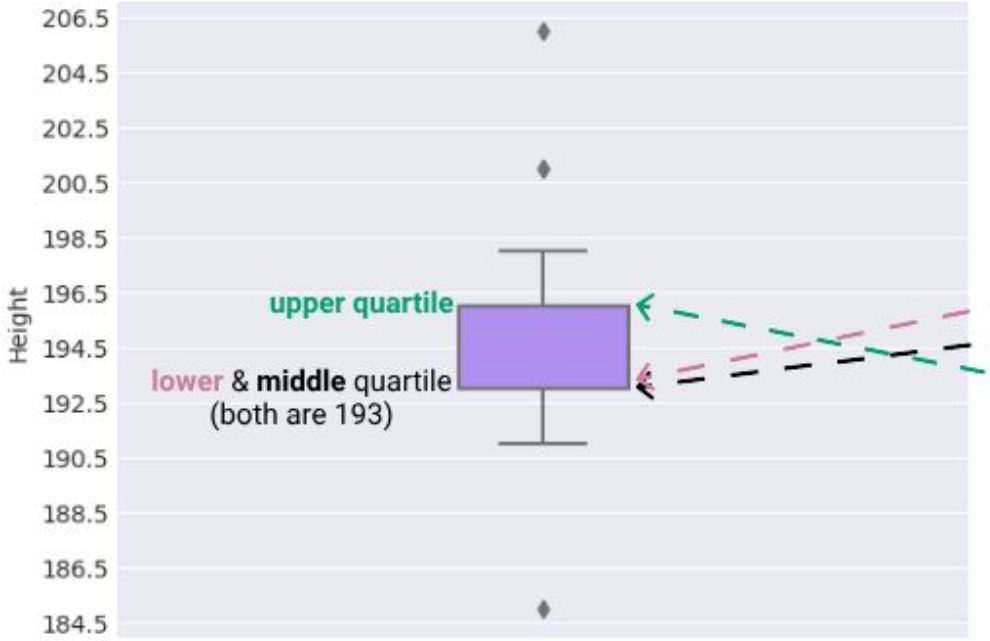


```
>> wnba[wnba['Pos'] == 'C']['Height'].describe()
count     25.000000
mean     194.920000
std        4.132392
min      185.000000
25%      193.000000
50%      193.000000
75%      196.000000
max      206.000000
Name: Height, dtype: float64
```

Outliers

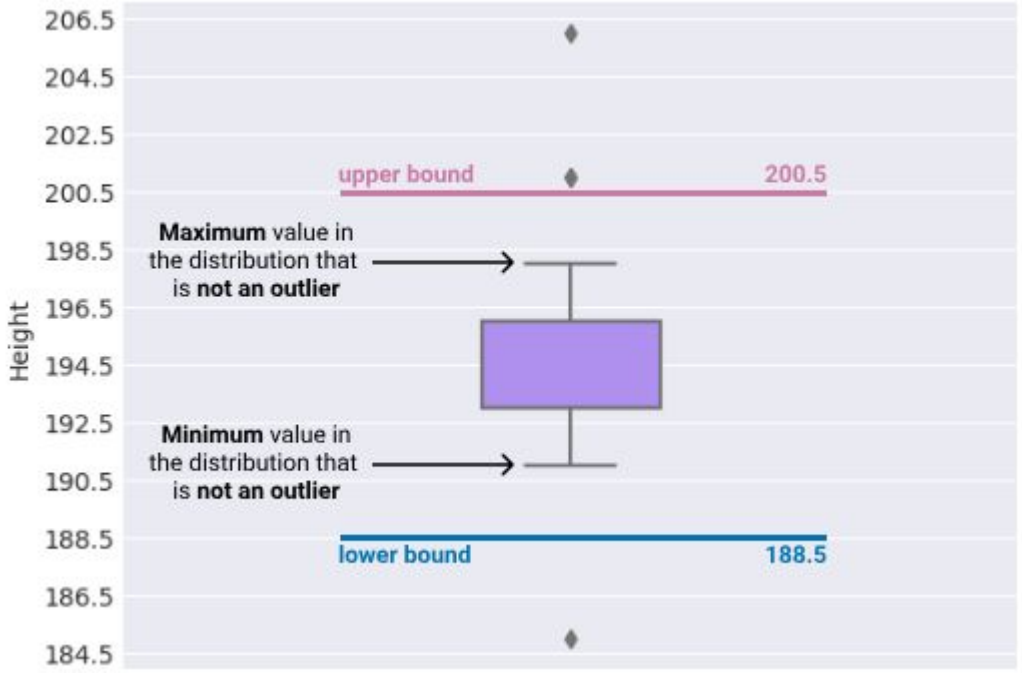
206.5
204.5
202.5
200.5
198.5
196.5
194.5
192.5
190.5
188.5
186.5
184.5
Height
upper bound
200.5
Maximum value in the distribution that is not an outlier
Minimum value in the distribution that is not an outlier
lower bound
188.5

Outliers

```
sns.boxplot(wnba[wnba['Pos'] == 'C']['Height'], whis = 1.5,
            orient = 'vertical', width = .45)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x1a180c4518>
```

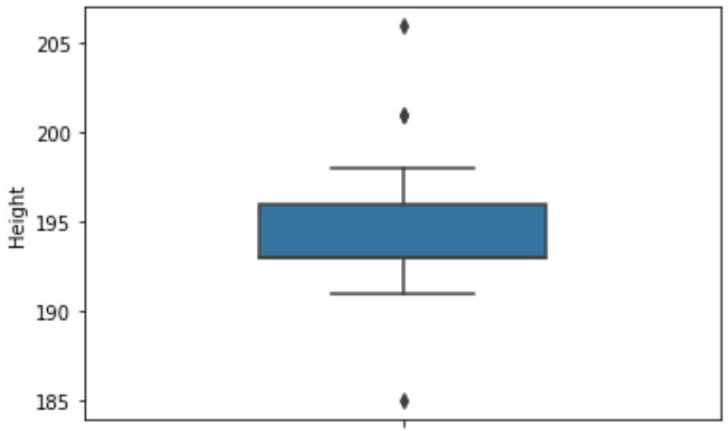


```
sns.boxplot(wnba[wnba['Pos'] == 'C']['Height'], whis = 4,
            orient = 'vertical', width = .45)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x1a18180208>
```

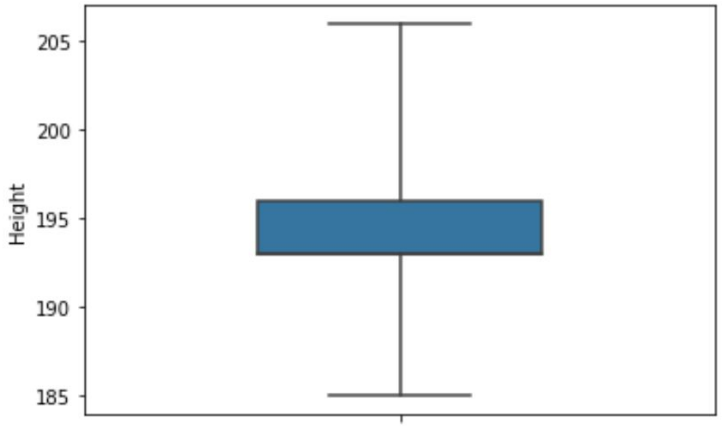

Scale of measurement	Graphs we can use to compare distributions
Nominal	
Ordinal	
Interval & Ratio	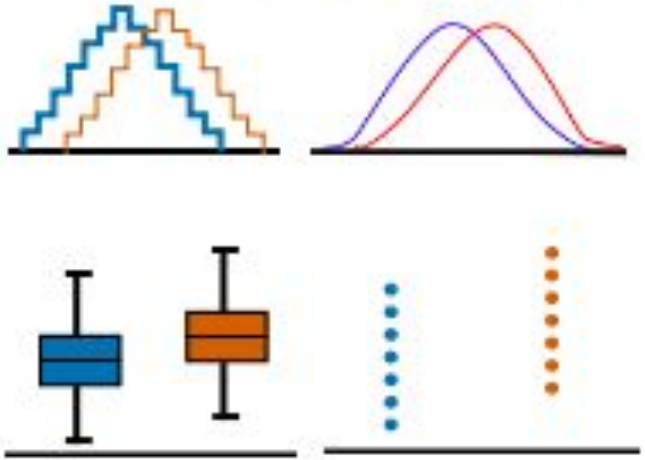

Data **source**

Getting **good** data

Id	Name	Salary	...	Gender
1	Mary Ann	$35 000	....	Female
2	Marc Downey	$55 000	....	Male
..	....	....	....	....
51	Juliet Ali	$45 000	....	Female
...	...	....	....	....
317	Jane Ace	$95 000	....	Female

Understand how the data is **structured** and **measured**

Organize the data in **comprehensible** forms to find patterns

Gender	Frequency
Male	147
Female	170

Visualize the patterns